

Advanced Research Computing: Practical session

1 Matrix multiplication

Write a program to accomplish the following:

- Take as user-input two arbitrary double-precision matrices from a text-file provided as a command-line argument.
- Multiply these two matrices together and output the solution to the terminal.
- You should not make use of pre-existing libraries (e.g. BLAS, Eigen, etc.) except for unit-testing.

For testing purposes, three files are provided in the `Practical` folder:

- `Matrix_Test_A_input.txt`: Two 3×3 matrices to be multiplied together.
- `Matrix_Test_B_bad_input.txt`: A 3×4 and a 3×3 matrix to be multiplied together. This clearly will not work.
- `Matrix_Test_C_large_input.txt`: A 64×96 and a 96×128 matrix to multiply together.

The file layout is as follows:

```
ROWS_1
COLS_1
A_00 A_01 A_02 A_03 ...
A_10 A_11 A_12 A_13 ...
...
A_N0 A_N1 ... A_NM
ROWS_2
COLS_2
B_00 B_01 B_02 B_03 ...
B_10 B_11 B_12 B_13 ...
...
B_N0 B_N1 ... B_NM
```

i.e. give the dimensions for each matrix, followed by the elements for each matrix. Numbers are separated using white-space.

1.1 Code running

In the standard case, your code would be run as:

```
./multiply ./Matrix_Test_A_input.txt
```

In order to save having to generate very large data-sets, your program should also accept matrix dimensions as options:

```
./multiply 1024 512 512 2048
```

which would generate two matrices of size 1024 512 and 512 2048, print them out, multiply them together, and then print their product.

See https://en.cppreference.com/w/cpp/numeric/random/uniform_int_distribution for how to generate reasonable random numbers. Try generating matrix elements in a small range including positive and negative number.

1.2 Optimization

You should make use of `valgrind` to try to optimize your code as far as possible. Points to note:

- Optimizations are only likely to become important for very large matrices, at least 100 100 if not 1000 1000 or larger.
- Be careful that your measured performance is not dominated by input/output.
- This level of optimization is difficult and involves a wide range of techniques, some of which are CPU and compiler dependent.
- As always, the optimal answer is an implementation of BLAS (or other external library).

1.3 Features

Your solution should have the following features:

- Make use of `CMake` to identify all compilers and other tools.
- Have all functions commented using Doxygen.
- Have unit-testing to ensure that invalid file-input (such as `Matrix_Test_B_bad_input.txt` above) is detected and does not result in seg-faults or other behaviour.

1.4 Suggestions

Some suggestions for how to lay out your code:

- Split your code into (at least) three files:
 - `Main.C`: Handles command-line parameters, and calls input/multiply/output routines.
 - `Multiply.C`: Contains functions to read input from a file, multiply two matrices, and output a result matrix.
 - `Test.C`: Contains unit-tests written with `Boost::Test` or `GoogleTest` that test the functionality of at least the input and multiply functions.
- Separate out the input/multiplication/output functionality into three separate functions within `Multiply.C`.
- Write tests that are designed to fail. For example, the invalid `Matrix_Test_B_bad_input.txt` above should be handled via an exception being thrown.
- The unit-tests could make use of the BLAS or Eigen libraries to check that the result is correct. (Or, use simple matrices as input and calculate the result by hand.)

1.5 Other hints

- Decide on your data layout for matrices first (maybe based on BLAS for compatibility).
- The BLAS routine for multiplying two double-precision matrices together is `cblas_dgemm`. An example program (for `float`) can be found at: <https://www.gnu.org/software/gsl/doc/html/cblas.html#examples>.
- Create trivial `structs` corresponding to different exceptions/error cases.
- Write skeleton-functions that are well-documented, i.e. detail the input/output conditions carefully.
- Write a simple `CMakeLists.txt` that initially just compiles the source-code.
- Extend `CMake` to compile two separate executables for normal running and for unit-tests.
- Add support for `clang-format` and `clang-tidy` that raises errors if code is not well-structured.